

## Summary

This code trains and evaluates a Random Forest Classifier to predict deforestation events based on land use and tree cover data. The input data consists of a stack of raster files, including land use plans, tree cover, and historical deforestation data. The model uses these raster files to predict deforestation events for the year 2012.

The input raster data is flattened and stacked into a single 2D array, `X_flat`. `NoData` values are removed from the input data (`X_cleaned`) and the target variable (`y_cleaned`) before splitting them into training and testing datasets.

The Random Forest Classifier is trained using the `X_train` and `y_train` datasets, and its performance is evaluated using cross-validation. The trained model is then used to predict deforestation events for the testing dataset (`X_test`). The model's performance is assessed using confusion matrices and classification reports for both the training and testing datasets.

Finally, the feature importances of the input variables (e.g., land use plans, tree cover) are calculated and visualized in a bar chart to understand the relative importance of each input variable in predicting deforestation events.

## Import Libraries and Constants

```
import os
import re
import sys
import numpy as np
import pandas as pd
import tempfile
import shutil
import matplotlib.pyplot as plt
import rasterio
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint as sp_randint

# Get the current working directory
current_dir = os.path.abspath('')
```

```

# Search for the 'constants.py' file starting from the current directory and moving up the
project_root = current_dir
while not os.path.isfile(os.path.join(project_root, 'constants.py')):
    project_root = os.path.dirname(project_root)

# Add the project root to the Python path
sys.path.append(project_root)

from constants import SERVER_PATH, OUTPUT_PATH, MASKED_RASTERS_DIR

#output- update this for subsequent runs
output_folder = os.path.join(OUTPUT_PATH[0], 'basic_rf_model')

# Where files for machine learning model should be located
# Directory containing the raster files
rasters_dir = MASKED_RASTERS_DIR[0]

```

## Create Stack

```

# helper function to read tiff files
def read_tiff_image(file_path):
    with rasterio.open(file_path) as src:
        return src.read(1)

# List of paths to the raster files excluding 'deforestation11_20_masked.tif'
feature_files = [os.path.join(rasters_dir, file_name) for file_name in os.listdir(rasters_dir)]

# Then you can use this list of raster_files to create feature_data_arrays and raster_data
feature_data_arrays = [read_tiff_image(file_path) for file_path in feature_files]
feature_data_flat = [data_array.flatten() for data_array in feature_data_arrays]

# Path to the y_file
y_file = os.path.join(rasters_dir, 'deforestation11_20_masked.tif')

```

```
feature_files
```

```
['/Users/romero61/../../../../capstone/pyforest/ml_data/output/masked_rasters/treecover_2010_masked.tif',  
 '/Users/romero61/../../../../capstone/pyforest/ml_data/output/masked_rasters/lup_10_masked.tif']
```

```
feature_files
```

```
['/Users/romero61/../../../../capstone/pyforest/ml_data/output/masked_rasters/treecover_2010_masked.tif',  
 '/Users/romero61/../../../../capstone/pyforest/ml_data/output/masked_rasters/lup_10_masked.tif']
```

```
# Find the dimensions of all the raster data arrays  
raster_shapes = [raster_data.shape for raster_data in feature_data_arrays]  
  
# Check if all raster data arrays have the same dimensions  
if len(set(raster_shapes)) > 1:  
    print("There are mismatching dimensions:")  
    for file_path, raster_shape in zip(raster_files, raster_shapes):  
        print(f"File: {file_path}, Shape: {raster_shape}")  
else:  
    print("All raster data arrays have the same dimensions.")  
    # Check the dimensions of all the raster data arrays  
    for i, data_array in enumerate(feature_data_arrays):  
        print(f"Raster {i}: {data_array.shape}")
```

All raster data arrays have the same dimensions.

Raster 0: (22512, 20381)

Raster 1: (22512, 20381)

## Stack and Flatten Data

```
# NoData Value  
no_data_value = -1  
  
# Stack the flattened raster data  
X_flat = np.column_stack(feature_data_flat)  
  
# Use the y_file obtained from the find_deforestation_file function
```

```

y = read_tiff_image(y_file).flatten()

# Remove rows with NoData values
'''checks each row in X_flat and creates a boolean array (valid_rows_X) that has the same
as the number of rows in X_flat. Each element in valid_rows_X is True if there is no NoData
the corresponding row of X_flat and False otherwise.'''
valid_rows_X = ~(X_flat == no_data_value).any(axis=1)

'''checks each element in the y array and creates a boolean array (valid_rows_y) that has
elements as y. Each element in valid_rows_y is True if the corresponding element in y is not
equal to the NoData value and False otherwise.'''
valid_rows_y = y != no_data_value

'''checks each element in the y array and creates a boolean array (valid_rows_y)
that has the same number of elements as y. Each element in valid_rows_y is True if the cor
in y is not equal to the NoData value and False otherwise.'''
valid_rows = valid_rows_X & valid_rows_y

'''creates a new array X_cleaned by selecting only the rows in X_flat that
correspond to the True elements in valid_rows.'''
X_cleaned = X_flat[valid_rows]

'''creates a new array y_cleaned by selecting only the elements in y that correspond
to the True elements in valid_rows.'''
y_cleaned = y[valid_rows]

```

To ensure your data cleaning steps have been applied correctly, you can check the following:

**NoData values have been removed:** You should confirm that there are no NoData values in your cleaned data. This can be done by asserting that there are no occurrences of `no_data_value` in `X_cleaned` and `y_cleaned`.

```

assert not (X_cleaned == no_data_value).any()
assert not (y_cleaned == no_data_value).any()

```

These assertions will throw an error if there is a NoData value in `X_cleaned` or `y_cleaned`

**Dimensions are correct:** The shapes of `X_cleaned` and `y_cleaned` should match along the row dimension (the first dimension for 2D array `X_cleaned` and the only dimension for 1D array `y_cleaned`).

```
print("Shape of X_cleaned:", X_cleaned.shape)
print("Shape of y_cleaned:", y_cleaned.shape)
```

Shape of X\_cleaned: (48181572, 2)

Shape of y\_cleaned: (48181572,)

Make sure the number of rows in X\_cleaned equals the number of elements in y\_cleaned.

**Confirm that the valid rows have been correctly identified:** You can do this by checking that the number of True values in valid\_rows equals the number of rows in X\_cleaned (or the number of elements in y\_cleaned).

```
assert valid_rows.sum() == X_cleaned.shape[0]
```

## Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X_cleaned, y_cleaned, test_size=0.2, r
```

## Class Balance Check

```
# Create pandas Series from your labels
y_train_series = pd.Series(y_train)
y_test_series = pd.Series(y_test)
y_cleaned_series = pd.Series(y_cleaned)

# Print balance of classes in training, testing, and whole dataset
print("Training data balance:\n", y_train_series.value_counts(normalize=True))
print("Testing data balance:\n", y_test_series.value_counts(normalize=True))
print("Whole dataset balance:\n", y_cleaned_series.value_counts(normalize=True))
```

Training data balance:

0 0.838179

1 0.161821

dtype: float64

Testing data balance:

0 0.838283

```

1    0.161717
dtype: float64
Whole dataset balance:
0    0.8382
1    0.1618
dtype: float64

```

The balance of your dataset seems to be roughly the same in both the training and testing sets, with about 83.8% of the instances belonging to class 0 (no deforestation) and 16.2% to class 1 (deforestation). This shows that the classes are quite imbalanced. Machine learning algorithms, including Random Forest, may have a bias towards the majority class in such situations, which could be one of the reasons why your model is not performing well on the minority class.

```

# Create a list to hold your feature file paths

# Define the labels for your features
feature_labels = ['TreeCover2010', 'LUP_10']

# Loop through your feature files
for feature_file, label in zip(feature_files, feature_labels):
    # Open the raster file
    with rasterio.open(feature_file) as src:
        # Read the raster data into a 2D array
        feature_data = src.read(1)

    # Calculate unique values and their counts
    unique_values, counts = np.unique(feature_data, return_counts=True)

    # Print the counts for each unique value
    for value, count in zip(unique_values, counts):
        print(f'{label} Value: {value}, Count: {count}')

    print('-----')

# Plot histogram
plt.figure(figsize=(10, 6))
plt.bar(unique_values, counts)
plt.title(f'{label} Distribution')
plt.xlabel('Values')
plt.ylabel('Counts')

```

```

plt.show()

for i, feature in enumerate(feature_labels):
    unique_values, counts = np.unique(X_cleaned[:, i], return_counts=True)

    # Print the counts for each unique value
    for value, count in zip(unique_values, counts):
        print(f'{feature} Value: {value}, Count: {count}')

    print('-----')

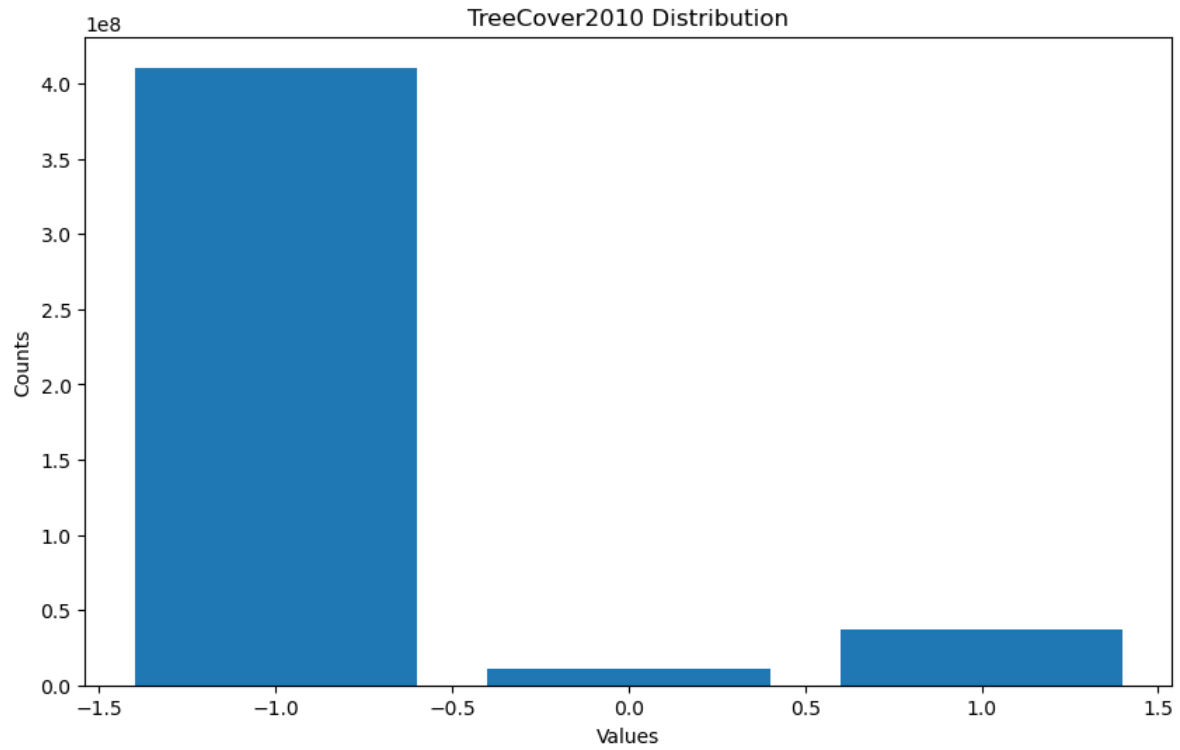
    # Plot histogram
    plt.figure(figsize=(10, 6))
    plt.hist(X_cleaned[:, i], bins=20)
    plt.title(f'{feature} Distribution')
    plt.xlabel('Values')
    plt.ylabel('Counts')
    plt.show()

```

```

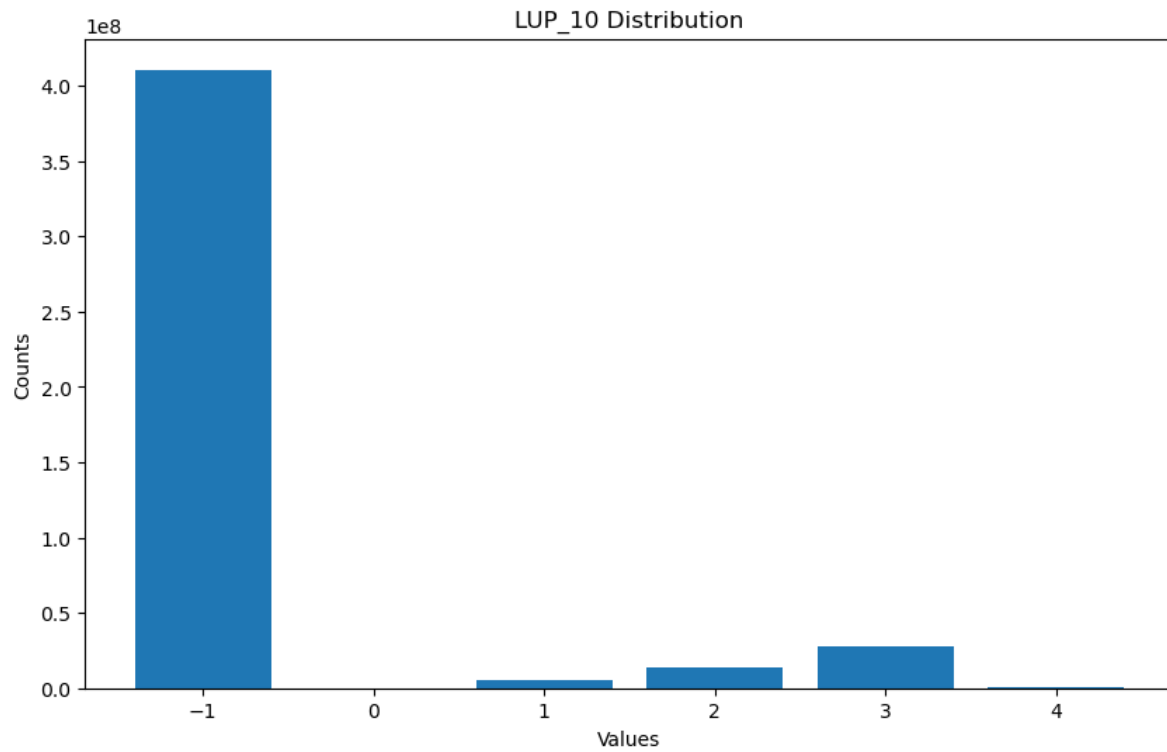
TreeCover2010 Value: -1, Count: 410635500
TreeCover2010 Value: 0, Count: 10878781
TreeCover2010 Value: 1, Count: 37302791
-----

```

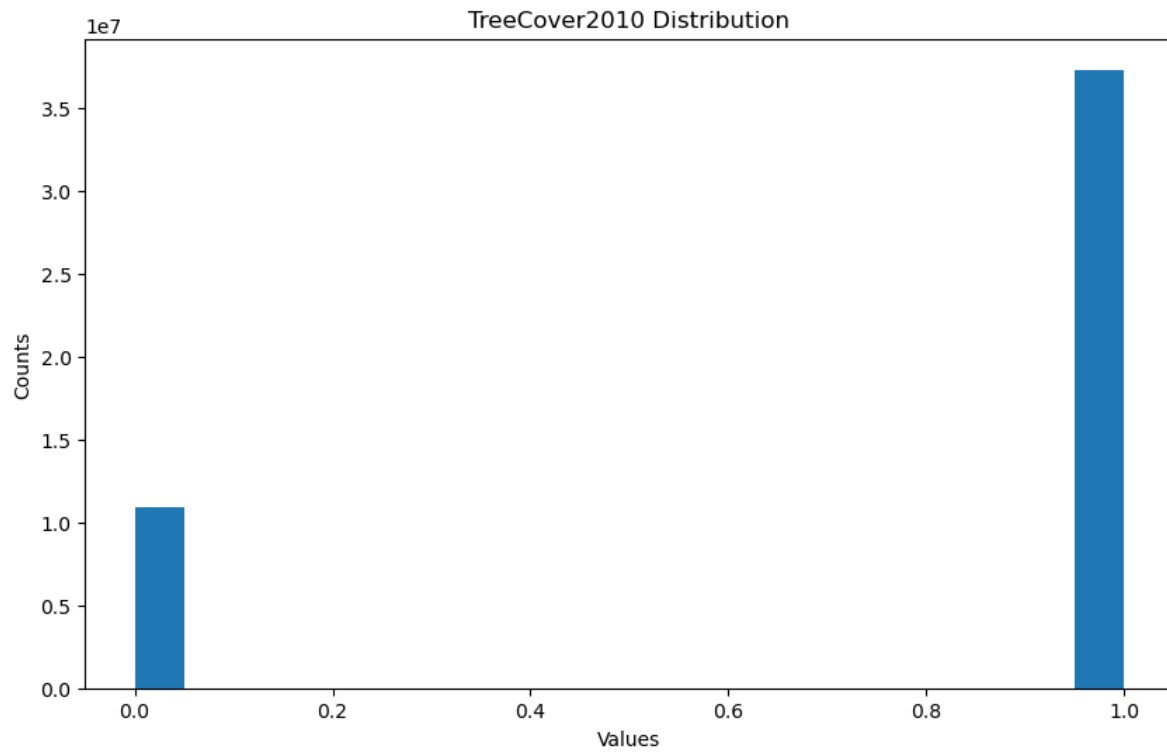


LUP\_10 Value: -1, Count: 410635500  
LUP\_10 Value: 1, Count: 5754071  
LUP\_10 Value: 2, Count: 13967435  
LUP\_10 Value: 3, Count: 27775793  
LUP\_10 Value: 4, Count: 684273  
-----

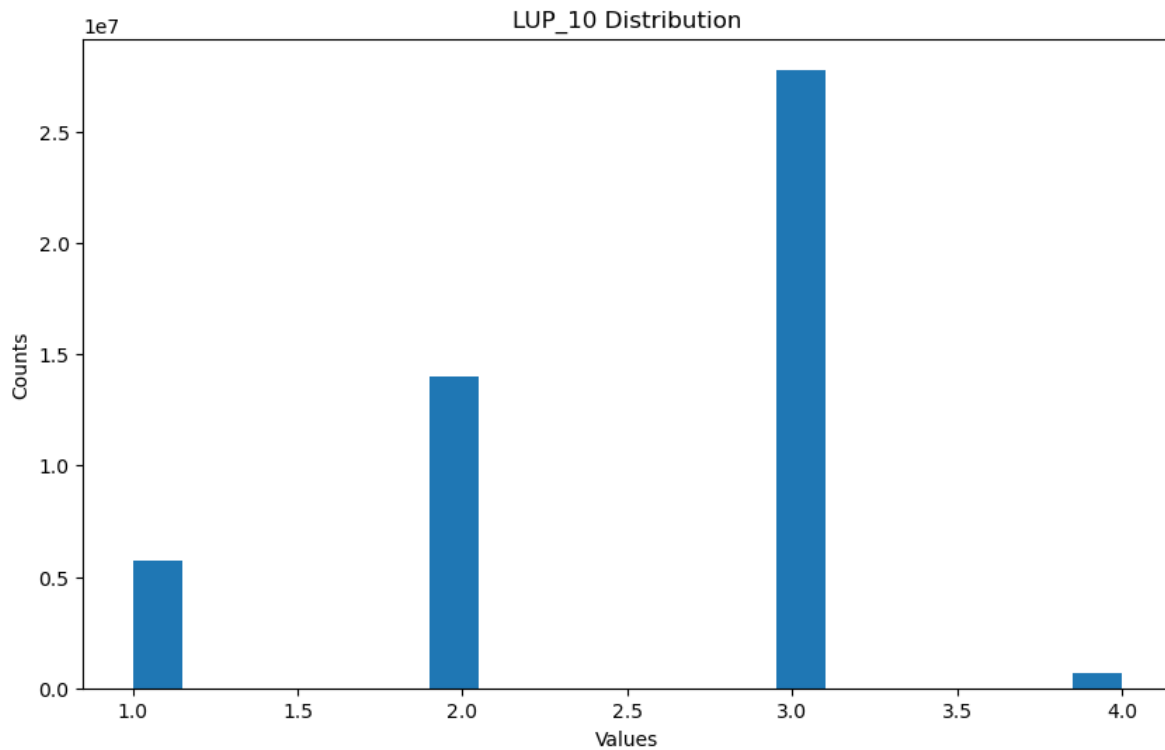




TreeCover2010 Value: 0, Count: 10878781  
TreeCover2010 Value: 1, Count: 37302791  
-----



LUP\_10 Value: 1, Count: 5754071  
LUP\_10 Value: 2, Count: 13967435  
LUP\_10 Value: 3, Count: 27775793  
LUP\_10 Value: 4, Count: 684273  
-----



## Simple Grid Search for a Random Forest model:

```
# Create a RandomForestClassifier instance
rfc = RandomForestClassifier(random_state=42, class_weight= 'balanced')

# Define a basic parameter grid
param_grid = {
    'n_estimators': [50, 100, 200],    # number of trees in the forest
    'max_depth': [None, 5, 10, 20]     # maximum depth of the tree
}

# Instantiate GridSearchCV
grid_search = GridSearchCV(rfc, param_grid, cv=5, n_jobs=30)
```

## Info on CV, fit, predict, predict\_proba

Cross-validation is a technique used to evaluate the performance of a machine learning model by training and testing it on different subsets of the dataset. It helps assess how well the model generalizes to new, unseen data and helps mitigate the risk of overfitting.

Here's how the 5-fold cross-validation works:

1. The entire dataset (`X_cleaned` and `y_cleaned`) is divided into 5 equally sized (or nearly equal) folds.
2. The model is trained and tested 5 times. In each iteration, one of the folds is used as the test set, and the remaining 4 folds are used to train the model.
3. For each iteration, the model's performance is evaluated using a chosen evaluation metric (in this case, accuracy, which is the default scoring method for `cross_val_score`).
4. Once all 5 iterations are completed, the performance scores are averaged to give a single cross-validation score.

By using cross-validation, you get a more reliable estimate of the model's performance because it's tested on different portions of the dataset. This helps to reduce the risk of overfitting and gives you a better understanding of how well your model generalizes to unseen data.

Cross-validation is performed before `clf.fit` to assess the performance of the model on the data without using the same data for both training and validation. It helps to understand how well the model is likely to generalize to new, unseen data before committing to training the final model.

If the cross-validation scores are satisfactory, you can proceed to train the final model using the entire dataset with `clf.fit`.

`clf.fit` is the method used to train the machine learning model on the provided dataset. In this case, it's training the Random Forest Classifier (denoted as `clf`) on the training dataset (`X_train` and `y_train`). The purpose of `clf.fit` is to learn the relationship between the input features (`X_train`) and the target variable (`y_train`) so that the model can make predictions on new, unseen data.

`clf.predict` is the method used to make predictions using the trained model. Once the model is trained with `clf.fit`, it can then be used to predict the target variable for new input features.

The model is predicting probabilities. The `RandomForestClassifier`, by default, outputs probabilities of class membership. It provides the probability of each pixel belonging to the deforested or non-deforested class. However, when you use `clf.predict()`, it returns the class with the highest probability, which is a **binary result (deforested or non-deforested)**.

The `clf.predict_proba()` function obtains probabilities instead of the binary result returning probabilities of each class. `y_proba = clf.predict_proba(X_cleaned)[: , 1]` extracts the probabilities of deforestation events (class 1) for all pixels.

```
# Fit to the training data
grid_search.fit(X_train, y_train)
```

```
GridSearchCV(cv=5,
             estimator=RandomForestClassifier(class_weight='balanced',
                                              random_state=42),
             n_jobs=30,
             param_grid={'max_depth': [None, 5, 10, 20],
                         'n_estimators': [50, 100, 200]})
```

## Examine Fit Results

`grid_search.best_params_` contains the hyperparameter combination that resulted in the highest average cross-validation score across the different folds during the grid search. This is useful information as it tells you which hyperparameters worked best for your model and data.

`grid_search.best_score_` is the highest mean cross-validation score achieved by the best hyperparameter combination found in the grid search. It gives you an idea of the model's performance with the optimal hyperparameters during the cross-validation process.

Best estimator: This provides the best estimator found by grid search. This is already fitted to the data and can be used for making predictions or further analysis.

CV Results: This provides a dictionary with various details about the grid search, like scores for each combination of parameters, time taken for fitting and scoring, etc. Note: `cv_results_` is a dictionary and can be quite verbose, you may want to convert it into a `DataFrame` for easier viewing.

Scorer: This provides the scoring function used by grid search.

Refit Time: This gives the time taken to refit the best estimator with the entire dataset.

```
# Print all available attributes and methods for the random_search object
all_attributes_methods = dir(grid_search)

# Filter out attributes and methods inherited from BaseSearchCV
specific_attributes_methods = [
    attribute for attribute in all_attributes_methods
```

```

        if attribute not in dir(GridSearchCV)
    ]

    print("Attributes and methods specific to GridSearchCV:")
    for attr in specific_attributes_methods:
        print(attr)

```

Attributes and methods specific to GridSearchCV:

```

best_estimator_
best_index_
best_params_
best_score_
cv
cv_results_
error_score
estimator
multimetric_
n_jobs
n_splits_
param_grid
pre_dispatch
refit
refit_time_
return_train_score
scorer_
scoring
verbose

```

```

grid_search.score

```

```

<bound method BaseSearchCV.score of GridSearchCV(cv=5,
          estimator=RandomForestClassifier(class_weight='balanced',
                                             random_state=42),
          n_jobs=30,
          param_grid={'max_depth': [None, 5, 10, 20],
                      'n_estimators': [50, 100, 200]})>

```

```

# Get the best parameters and the corresponding score

```

```

best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best parameters:", best_params)
print("Best cross-validation score:", best_score)

```

Best parameters: {'max\_depth': None, 'n\_estimators': 50}  
 Best cross-validation score: 0.5858522619204243

After fitting the GridSearchCV, you can evaluate the performance of the best model on the test data (X\_test and y\_test) using the best\_estimator\_ attribute of the grid\_search object:

```

best_estimator = grid_search.best_estimator_

cv_results = grid_search.cv_results_

cv_results_df = pd.DataFrame(grid_search.cv_results_)

scorer = grid_search.scorer_

refit_time = grid_search.refit_time_

print("Best estimator:", best_estimator)
print("CV Results:", cv_results_df)
print("Scorer function:", scorer)
print("Refit time (seconds):", refit_time)

```

Best estimator: RandomForestClassifier(class\_weight='balanced', n\_estimators=50, random\_state=42)

CV Results:	mean_fit_time	std_fit_time	mean_score_time	std_score_time	\
0	546.787241	58.240538	21.566601	3.236626	
1	1017.598757	75.864049	46.208107	9.260437	
2	1918.190541	115.549766	73.645717	1.699839	
3	474.310703	9.245638	21.489198	3.174777	
4	1070.831887	49.192505	53.257504	7.381053	
5	1934.186380	90.319932	73.316022	1.469383	
6	466.367223	42.144907	22.921606	3.658064	
7	1021.736046	90.103013	45.873498	6.964352	
8	1686.582774	28.455681	70.935093	0.173656	

9	534.640818	40.057919	24.653716	4.052579
10	904.753022	26.085620	36.163089	0.194958
11	1552.772714	32.792583	69.884495	1.043727

	param_max_depth	param_n_estimators	\
0	None	50	
1	None	100	
2	None	200	
3	5	50	
4	5	100	
5	5	200	
6	10	50	
7	10	100	
8	10	200	
9	20	50	
10	20	100	
11	20	200	

	params	split0_test_score	\
0	{'max_depth': None, 'n_estimators': 50}	0.585995	
1	{'max_depth': None, 'n_estimators': 100}	0.585995	
2	{'max_depth': None, 'n_estimators': 200}	0.585995	
3	{'max_depth': 5, 'n_estimators': 50}	0.585995	
4	{'max_depth': 5, 'n_estimators': 100}	0.585995	
5	{'max_depth': 5, 'n_estimators': 200}	0.585995	
6	{'max_depth': 10, 'n_estimators': 50}	0.585995	
7	{'max_depth': 10, 'n_estimators': 100}	0.585995	
8	{'max_depth': 10, 'n_estimators': 200}	0.585995	
9	{'max_depth': 20, 'n_estimators': 50}	0.585995	
10	{'max_depth': 20, 'n_estimators': 100}	0.585995	
11	{'max_depth': 20, 'n_estimators': 200}	0.585995	

	split1_test_score	split2_test_score	split3_test_score	\
0	0.585956	0.586016	0.585466	
1	0.585956	0.586016	0.585466	
2	0.585956	0.586016	0.585466	
3	0.585956	0.586016	0.585466	
4	0.585956	0.586016	0.585466	
5	0.585956	0.586016	0.585466	
6	0.585956	0.586016	0.585466	
7	0.585956	0.586016	0.585466	
8	0.585956	0.586016	0.585466	
9	0.585956	0.586016	0.585466	



10	0.585956	0.586016	0.585466
11	0.585956	0.586016	0.585466

	split4_test_score	mean_test_score	std_test_score	rank_test_score
0	0.585829	0.585852	0.000204	1
1	0.585829	0.585852	0.000204	1
2	0.585829	0.585852	0.000204	1
3	0.585829	0.585852	0.000204	1
4	0.585829	0.585852	0.000204	1
5	0.585829	0.585852	0.000204	1
6	0.585829	0.585852	0.000204	1
7	0.585829	0.585852	0.000204	1
8	0.585829	0.585852	0.000204	1
9	0.585829	0.585852	0.000204	1
10	0.585829	0.585852	0.000204	1
11	0.585829	0.585852	0.000204	1

Scorer function: <function \_passthrough\_scorer at 0x7f2b78a31300>

Refit time (seconds): 459.01079297065735

## Evaluate the model performance using your preferred metrics

e.g., confusion matrix, classification report, accuracy, F1-score, etc.

```
best_model = grid_search.best_estimator_
```

```
# Predictions for test data
y_pred = best_model.predict(X_test)
```

Evaluate the performance of your model by comparing the predicted labels (`y_pred`) with the true labels (`y_test`). You can use various metrics such as confusion matrix, classification report, accuracy, F1-score, etc.:

```
from sklearn.metrics import accuracy_score, f1_score, classification_report

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Calculate F1-score (use 'weighted' or 'macro' depending on your problem)
f1 = f1_score(y_test, y_pred, average='weighted')
```

```

print("F1-score:", f1)

# Print classification report
report = classification_report(y_test, y_pred)
print("Classification report:\n", report)

```

Accuracy: 0.5859063345272545

F1-score: 0.6391117670833032

Classification report:

	precision	recall	f1-score	support
0	0.93	0.55	0.69	8077959
1	0.25	0.79	0.38	1558356
accuracy			0.59	9636315
macro avg	0.59	0.67	0.53	9636315
weighted avg	0.82	0.59	0.64	9636315

## Confusion Matrix

```

# Predictions for train data
y_pred_train = best_model.predict(X_train)

# Confusion matrix and classification report for train data
train_cm = confusion_matrix(y_train, y_pred_train)
train_cr = classification_report(y_train, y_pred_train)
print("Training confusion matrix:")
print(train_cm)
print("Training classification report:")
print(train_cr)

```

Training confusion matrix:

[[17670601 14637219]

[ 1326212 4911225]]

Training classification report:

	precision	recall	f1-score	support
0	0.93	0.55	0.69	32307820
1	0.25	0.79	0.38	6237437

accuracy			0.59	38545257
macro avg	0.59	0.67	0.53	38545257
weighted avg	0.82	0.59	0.64	38545257

```
# Calculate feature importances and the standard deviation for those importances
importances = best_model.feature_importances_
std = np.std([tree.feature_importances_ for tree in best_model.estimators_], axis=0)

# list of feature names corresponding to the input bands of your raster stack
feature_names = ['TreeCover2010', 'LUP']
# Create a sorted list of tuples containing feature names and their importances:
sorted_features = sorted(zip(feature_names, importances, std), key=lambda x: x[1], reverse=True)

# Create a bar chart
fig, ax = plt.subplots()

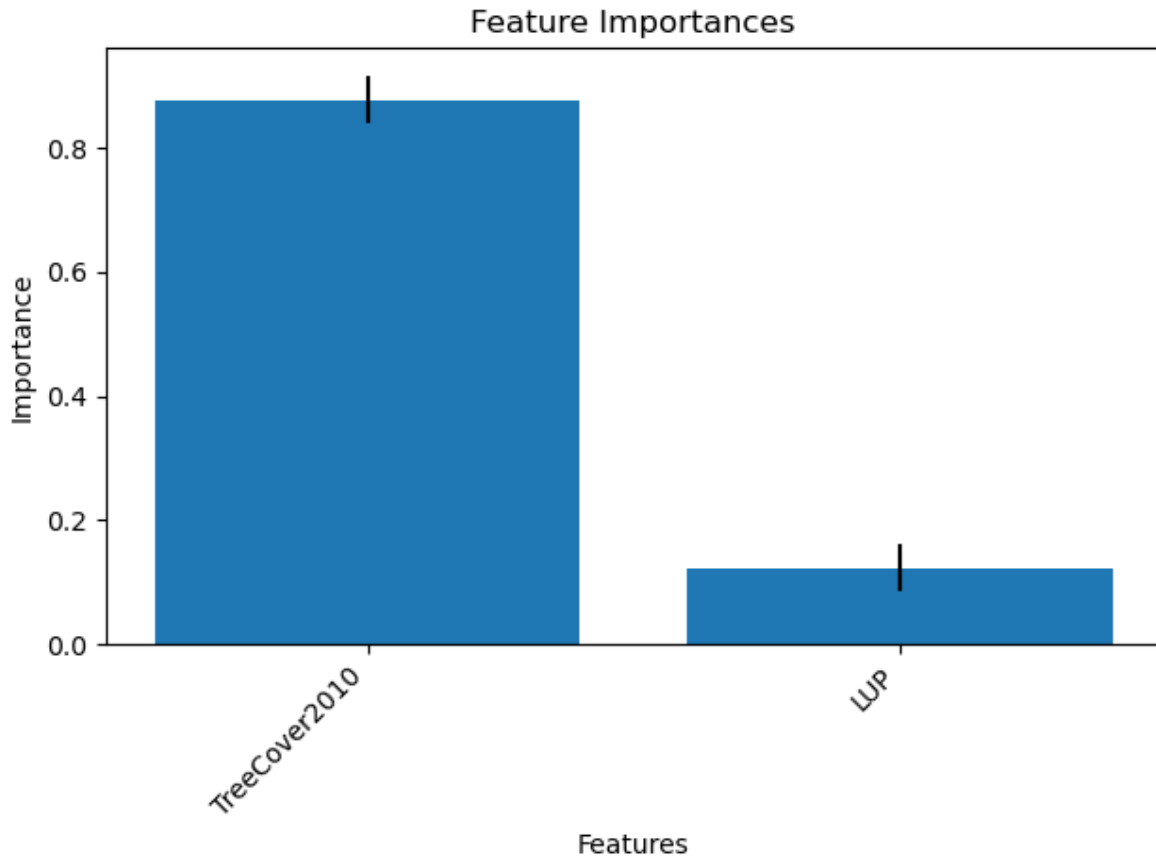
# Set the feature names as x-axis labels
ax.set_xticklabels([item[0] for item in sorted_features], rotation=45, ha='right')
ax.set_xticks(range(len(sorted_features)))

# Set the y-axis labels as importances
ax.bar(range(len(sorted_features)), [item[1] for item in sorted_features], yerr=[item[2] for item in sorted_features])

# Set the title and labels for the chart
ax.set_title('Feature Importances')
ax.set_xlabel('Features')
ax.set_ylabel('Importance')

# Display the chart
plt.tight_layout()
plt.show()
```

```
/tmp/ipykernel_3700312/2283730908.py:15: UserWarning: FixedFormatter should only be used to
ax.set_xticklabels([item[0] for item in sorted_features], rotation=45, ha='right')
```



## Probabilities for deforestation

```
# Predict probabilities for deforestation events
y_proba = best_model.predict_proba(X_cleaned)[: , 1]

# Predicts the
# Create a probability raster by filling in the valid pixel values
prob_raster = np.full(y.shape, no_data_value, dtype=np.float32)
prob_raster[valid_rows] = y_proba
prob_raster = prob_raster.reshape(feature_data_arrays[0].shape)

# Save the probability raster as a GeoTIFF file
if not os.path.exists(output_folder):
    os.makedirs(output_folder)
```

```

output_file = os.path.join(output_folder, "deforestation_prob.tiff")

with rasterio.open(y_file) as src:
    profile = src.profile
    profile.update(dtype=rasterio.float32, count=1)

prob_raster_resaped = prob_raster.reshape((1, prob_raster.shape[0], prob_raster.shape[1]))

with rasterio.open(output_file, 'w', **profile) as dst:
    dst.write_band(1, prob_raster_resaped[0])

```

## Tuning Strategies

```

# Complex Grid
'''# Set the range of values for each hyperparameter
param_grid = {
    'n_estimators': [100, 200],
    'criterion': ['gini', 'entropy'],
    'max_depth': [10, 20, 30, 40],
    'min_samples_split': [ 5, 10, 20],
    'min_samples_leaf': [2, 5, 10],
    'max_features': ['sqrt', None],
    'bootstrap': [True],
    'class_weight': ['balanced']
}

# Instantiate the RandomForestClassifier
clf = RandomForestClassifier(random_state=0)

# Set up the GridSearchCV
grid_search = GridSearchCV(clf, param_grid=param_grid, cv=5, n_jobs=19)'''

# Randomized Search
# Set the range of values for each hyperparameter
'''param_dist = {
    "n_estimators": sp_randint(100, 300),

```

```

        'criterion': ['gini'],
        'max_features': ['sqrt', None],
        "max_depth": sp_randint(1, 20),
        "min_samples_split": sp_randint(2, 11),
        "min_samples_leaf": sp_randint(1, 11),
        "bootstrap": [True],
        'class_weight': ['balanced']
    }

    # Instantiate the RandomForestClassifier
    clf = RandomForestClassifier(random_state=0)

    # Set up the RandomizedSearchCV
    random_search = RandomizedSearchCV(
        clf, param_distributions=param_dist, n_iter=20, cv=5, random_state=0, n_jobs=19
    )

```